



# Empirical Study of Partial Evaluation of Matrix and String Algorithms

Ilya Balashov  
Daniil Berezun  
Semyon Grigorev

Saint Petersburg State University

SYRCoSE'21

# Introduction: Partial Evaluation and Huge Programs

- Huge programs often have core parts, which are heavily loaded or employed by a large number of other parts
  - ▶ Matrix multiplication algorithms in BLAS
  - ▶ Pattern and automaton matching in programs on strings
- Optimization of core parts  $\equiv$  optimization of the whole program
- It is difficult to write well-optimized programs
  - ▶ Helper tools and methods are needed
- *Partial Evaluation* — one of such a methods

# In This Report...

Partial evaluation applied to a linear algebra and string algorithms.

- 1 Brief introduction to the background
  - ▶ Partial evaluation
  - ▶ Algorithms for the experiments
- 2 Algorithms implementation
- 3 Experimental design
- 4 Results
- 5 Related & future work

# Background: Partial Evaluation

- Method of program transformation
  - ▶ Tool — Partial Evaluator
- $\llbracket F \rrbracket[a, b] = \llbracket F_b \rrbracket[a]$ 
  - ▶  $F_b$  could be faster, smaller and easier than  $F$
- Also used for compiler generation
  - ▶ Futamura projections
- We chose AnyDSL framework
  - ▶ DSL Impala and Artic
  - ▶ Previously applied in several areas:
    - ★ Image processing
    - ★ Bioinformatics
    - ★ Ray tracing

# Background: Algorithms

- Matrix algorithms
  - ▶ Matrix multiplication
  - ▶ Kronecker (tensor) product
- String processing algorithms
  - ▶ Pattern matching
  - ▶ Regular expression matching (in matrix form)
- Applied widely both in science and industry
  - ▶ GraphBLAS primitives
  - ▶ KMP-test
  - ▶ Utilities like Grep

# Experimental Design: Questions

Research questions:


- Q1 Does partial evaluated benefits string and matrix-based graph algorithms performance comparing to their basic versions?
- Q2 In which degree partially evaluated algorithms code performance gets closer to their state-of-art implementations?

# Experimental Design: Datasets


- For matrix algorithms
  - ▶ SuiteSparse matrix collection
  - ▶ Harwell-Boeing matrix collection
- For algorithms on string
  - ▶ Autogenerated strings
  - ▶ Traffic dumps
  - ▶ Regular expressions from `regexlib.com` catalog
- We chose data with the diverse configuration
  - ▶ Tried to cover more basic test cases
  - ▶ Used degenerate cases
  - ▶ A lot of different  $\implies$  less threats to validity

# Results: Matrices

- Intel i5-7440HQ, 16GB RAM
- (Non)Specialized code in AnyDSL Impala
- Google Benchmark
- Only “interesting” cases are shown




Matrix multiplication





Kronecker product



# Results: Strings




Pattern matching




Regular expressions

# Comparison with SuiteSparse GraphBLAS




Matrix multiplication




Kronecker product

# Comparison with (e)Grep



Pattern matching



Regular expressions

# Results: In General

- 10% to 100 times speedup for nearly all tested algorithms comparing to non-partially evaluated code
- 2–5 time to 3 orders win by time in comparison with (e)Grep
- SuiteSparse won 5 times against partially evaluated code
  - ▶ Good result for an semi-automatic optimization
  - ▶ We still won 10+ times comparing to non-evaluated code
  - ▶ The code was optimized on compile-time — no need for a heavy library

# Related and Future Work

- Wide set of directions for the future
  - ▶ More complex algorithms
  - ▶ Another tools & languages (AnyDSL Artic)
  - ▶ GPGPU
- Related covers a lot of topics
  - ▶ Ray tracing/Bioinformatics/Image processing (AnyDSL team)
  - ▶ Viterbi algorithm specialization (Ivan Tyulyandin, SEIM'21)
  - ▶ CUDA specialization (Alexey Tyurin, PPOPP'20)

## Appendix: Full Matrix List

	Size	Not Null	Symmetry, %	Value Type
<i>bcsstk16</i>	4884	147631	100	real
<i>fs_183_1</i>	183	1069	41.8	real
<i>can_256</i>	256	2916	100	binary
<i>eye3</i>	3	3	100	binary
<i>2blocks</i>	4	8	100	binary
<i>cover</i>	8	12	16.67	binary
<i>mycielskian3</i>	6	5	0	binary
<i>trec5</i>	8	12	0	real

Table: Matrices used in partial evaluation experiments

## Appendix: Full Results of Matrix Multiplication Experiments

Time, ns. Spec/NoSpec/SP $\Delta < 0.01\%$	$\times$ <i>eye3</i>	$\times$ <i>2blocks</i>	$\times$ <i>cover</i>	$\times$ <i>my...3</i>	$\times$ <i>trec5</i>
<i>bcsstk16</i> $\times$	93608 121855 2270	133434 157850 7064	364772 4842889 8559	171085 2129094 511	308535 5226893 505
<i>fs_183_1</i> $\times$	7796 6752 2553	20187 42353 12310	6928 38250 9796	1358 15194 506	6078 42493 507
<i>can_256</i> $\times$	1016 1177 2259	5106 38221 6549	20339 66987 9409	2561 23105 503	9548 62668 506

Table: Execution times for matrix multiplication experiments

## Appendix: Full Results of Kronecker Product Experiments

Time, ns. Spec/NoSpec/SP $\Delta < 0.01\%$	$\otimes$ <i>eye3</i>	$\otimes$ <i>2blocks</i>	$\otimes$ <i>cover</i>	$\otimes$ <i>my...3</i>	$\otimes$ <i>trec5</i>
<i>bcsstk16</i> $\otimes$	140628 140744 901878	276222 3032308 2145104	433397 4307538 4420688	276433 1967189 2958016	481805 4571625 1440326
<i>fs_183_1</i> $\otimes$	916 934 25833	2186 21272 45159	3046 31732 88847	1838 14533 35109	3146 34356 47912
<i>can_256</i> $\otimes$	1159 1069 35162	2772 30841 60600	4512 45731 130084	2736 22079 43479	4576 49512 61500

Table: Execution times for Kronecker product experiments



# Appendix: Pure C Comparison

